

JSC270 Winter 2022 Assignment 4

Natural Language Processing (NLP)

Submission Deadline: Friday, April 8th at 10:59AM EST

Presentation Dates: Monday, April 4th, between 12:00PM and 1:00PM EST (during class) and Wednesday, April 6th, between 12:00PM and 2:00PM EST (during lab)

What to submit:

1. A **PDF report** containing your written responses to Parts I and II. Include any figures and tables that are asked for in the assignment or are necessary to understand your analysis.
 - In the beginning of the report, please specify the breakdown of your work indicating the contributions of each partner. For example, if all the work was done jointly or if one of you did most of the analysis, and another was mostly responsible for the presentation and writing the report.
2. A Google Colab **notebook** containing relevant python code. Provide a link to this notebook at the top of your PDF report. Make sure you change the share settings (top right of the colab you'll see a "Share" icon) to "Anyone with a Link" and "Commenter."
 - Alternatively, you may put the code on GitHub, and instead provide a link to that repo following the instructions provided in Part I of Assignment 2.
 - Use section headers in your notebook to indicate which pieces of code correspond to which question.
3. Your **custom Twitter data** for Part II (in CSV format).
4. **Slides** for the presentation of Part III (in PDF format).

You only need to turn in **one PDF report, notebook, slides, and data per group**. Make sure your name and your partner's name are at the top of the PDF and the notebook.

Where to submit: The PDF report, notebook, slides, and data are to be submitted through Quercus by the due date. Submit the files under your name or your partner's name (not both!).

Grading Scheme: This assignment is worth a total of 60 marks (5 marks for signing up for groups by 10:59 am on 3/14). *This assignment is to be completed with your partner.*

Suggestions:

- Start early! You want to make sure you have enough time.
- For Part II, obtaining approval for use of the Twitter API takes around 2 days. Please plan accordingly when completing the assignment. Instructions were provided in Lab 7.

Part I: Sentiment Analysis with a Twitter Dataset (20 pts)

For this question, you'll be using a Twitter dataset containing just under 45,000 tweets related to COVID-19. These data come from a fairly recent Kaggle competition, which you can read more about [here](#). To help with modeling, we've simplified the data in the following ways:

- The tweets are split into training (covid-tweets-train.csv) and test sets (covid-tweets-test.csv). The training set contains 41155 observations, while the test set contains 3798. This corresponds to (roughly) a 91.5/8.5 split. These files are available in the assignments section of the course website.
- The dataset contains the tweets (as strings) and labels $y_i \in \{0,1,2\}$, corresponding to sentiment {negative, neutral, positive}. You may remove observations with missing labels or sentiment scores that are strings.

In this question, you'll be training classifiers to predict whether the tweet is positive, negative, or neutral, based only on the tweet itself. This is called sentiment analysis, and it is a common ML problem. To give you a sense of how sentiment is labeled, below are tweets from each of the three classes:

- **Positive:** *Hey guysss! We have a very important message to all our friends, family amp; viewers out there. Be aware of Corona Virus and follow these procedures to help keep you and your family safe.#aadyasitara? #aadya #sitara #coronavirus #coronaalert covid_19 #staysafe #sanitizer #covid*
<https://t.co/pHWGsygy06>
- **Neutral:** *The top food items to stock up on in case you are quarantined #Coronavirus home stays are on people's minds. Here are ten foods for stocking a pantry to support physical and mental health.* <https://t.co/jrvxRlb57L>
- **Negative:** *Just thinking if we have to close schools, how many children may lose access to main source of meals. And what will happen to demand at food banks.*
#trusselltrust #coronavirus #foodbanks #poverty #uk <https://t.co/Mw71RzC6z1>

Note: Each of the preprocessing steps that follow should be applied to both the training and test sets. You may find it useful to write these steps as functions that can be applied to both datasets (and potentially even used for Part II). None of these preprocessing steps should require more than a few lines of code.

A) (1 pt) Consider the training data. What is the balance between the three classes? In other words, what proportion of the observations (in the training set) belong to each class?

B) (1 pt) Tokenize the tweets. In other words, for each observation, convert the tweet from a single string of running text into a list of individual tokens (possibly with punctuation), splitting on whitespace. The result should be that each observation (tweet) is a list of individual tokens.

C) (1 pt) Using a regular expression, remove any URL tokens from each of the observations.

Hint: In this dataset, all such tokens begin with “http”.

D) (2 pts) Remove all punctuation (.,?!;:'") and special characters(@, #, +, &, =, \$, etc). Also, convert all tokens to lowercase only. Can you think of a scenario when you might want to keep some forms of punctuation?

E) (1pt) Now stem your tokens. This will have the effect of converting similar word forms into identical tokens (e.g. run, runs, running → run). Please specify which stemmer you use.

Note: There are several different stemmers available through nltk and Scikit-learn. I recommend the Porter stemmer, but you may use a different one if you wish.

F) (1pt) Lastly, remove stopwords. Using the english stopwords list from nltk, remove these common words from your observations. This list is very long (I think almost 200 words), so remove only the first 100 stopwords in the list.

G) (2 pts) Now convert your lists of words into vectors of word counts. You may find Scikit-learn’s CountVectorizer useful here. What is the length of your vocabulary?

Hint: The matrix of counts will be $D \times V$, where D is the number of documents (tweets), and V is the number of features (word counts).

H) (4 pts) Recall the definition of the Naive Bayes model. If each document (tweet) is a collection of words (w_1, \dots, w_N) belonging to class C_k ($k = 0, 1, 2$), then the Naive Bayes approach models the probability of each tweet belonging to class k :

$$\begin{aligned} P(C_k | w_1, \dots, w_N) &\propto P(w_1, \dots, w_N | C_k) P(C_k) \\ &= P(C_k) \prod_{i=1}^N P(w_i | C_k) \end{aligned}$$

The last equality follows from our “naive” assumption that words are conditionally independent given class. The probabilities are estimated using the frequencies of words within each class (bag of words), and we assign the class label according to which of the 3 posterior class probabilities ($P(C_k | w_1, \dots, w_N)$) is the highest.

Fit a Naive Bayes model to your data. Report the training and test error of the model. Use accuracy as the error metric. Also, report the 5 most probable words in each class, along with their counts. You might find Scikit-learn’s MultinomialNB() transformer useful. Use Laplace smoothing to prevent probabilities of zero.

I) (2 pts) Would it be appropriate to fit an ROC curve in this scenario? If yes, explain why. If no, explain why not.

J) (2 pts) Redo parts G-H using TF-IDF vectors instead of count vectors. You might find Scikitlearn's `TfidfVectorizer()` transformer useful. Report the training and test accuracy. How does this compare to the accuracy using count vectors?

K) (3 pts) Recall lemmatization converts each word to its base form, which is a bit stronger than simply taking the stem. Redo parts E-H using TF-IDF vectors instead of count vectors. This time use lemmatization instead of stemming. Report train and test accuracy. How does the accuracy with lemmatization compare to the accuracy with stemming?

Note: Like stemmers, there are multiple lemmatizers you might use. We recommend the WordNet lemmatizer offered by nltk.

Bonus (1 pt): Is the Naive Bayes model generative or discriminative? Explain your response.

Part II: Having fun with NLP using the Twitter API (20 pts)

For this question, you'll extract your own Twitter data using Twitter's python API. The free version of this interface contains several different search parameters, but only allows for extraction of tweets from the past seven days (exactly). This means that each group's dataset will be unique.

You'll devise a research question that can be answered using your custom data (and possibly additional data if you require it). You'll then answer that question using a machine learning model of your choice. The analysis should be guided by the five components listed later in the section.

The Python Twitter API

The python wrapper for Twitter's API is called tweepy. The standard(free) version allows you to extract tweets from up to 7 days ago, with a limit of 18,000 tweets in a 15 minute window.

In order to use this package, you'll first need to obtain some credentials from Twitter. Following the instructions from [Lab 7](#) to obtain credentials. Make sure to record:

- A Consumer API Key
- A Consumer API Secret Key
- An Access Token
- An Access Token Secret

We'll use it to connect to Twitter through tweepy. Note that the developer dashboard will only show you these 4 pieces of information once. For security, if you ever want to view them again, Twitter will generate new ones for you.

Analysis Structure

Although each group will have a different analysis, the general structure should look the same. Specifically, you should cover each of the 5 components described below in your analysis and report. Include a section in your report for each component. The report should be **no longer than 5 pages** (double spaced, 12 pt font, 1 inch margins).

(I) **Problem description and motivation.** Describe the question you will answer. Why is this question difficult to solve? How will your dataset help solve it? Is there any literature related to existing approaches? Why will your analysis be different? Make sure you provide proper citation of any other approaches you reference in your analysis.

Note: It's okay if another analysis has taken an approach similar to yours. You'll have a strong baseline comparison for your results.

(II) **Describe the data.** Describe the data you've extracted. What exactly were the extraction parameters you used in the API, and why? State the number of observations and the number of features. Have other similar data been analyzed, and if so, how? Did other similar works of research generate conclusions similar to yours? What are the limitations of your data? What are the strengths of your data?

Note: If, due to computational constraints, you can only carry out your analysis using a subset of the data you extract, that's okay. Just be clear about the data you selected and why.

(III) **Exploratory data analysis.** Summarize your findings from EDA, including:

- Computing the basic summary statistics of your data.
- Making basic plots to help you better understand the data.
- Performing any necessary preprocessing steps to ensure your data is ready for model-fitting. Many of these steps will be similar to what was done in Part I (that's okay!). You might generate new features based on the raw text.
- Reporting any interesting findings you come across.

Your EDA is a good opportunity to generate some visualizations for your presentation (Part III).

(IV) **Describe your machine learning model.** Identify the machine learning task you'll be using to answer your question (eg. binary classification, multiclass-classification,

clustering - we will get to this topic). Explain how the model works. Why is this model an appropriate choice for your research problem? Is this a supervised or unsupervised model? What are its strengths and weaknesses? How will you evaluate the model's performance? Are there any baseline models to which you can compare your results? You can use any of the models discussed in class, but you are not constrained to just those (some additional examples will be listed later in this section).

Note: For supervised learning models, you will probably have to generate the labels yourself.

(V) **Results and Conclusions.** Describe the results of your analysis. How well does your model perform? Did you run into any issues (e.g. overfitting), and if so, how did you manage them? What is the significance of your results? Do your model parameters have a meaningful interpretation? Why did your model work better than existing approaches (or if it failed, why did it fail)? If you had more time, or could collect more data and what would you do differently?

Note: It is ok if your analysis does not turn out the way you want - this happens all the time! Be clear about your thought process in doing the analysis, what might explain what went wrong, and what could be done to improve results.

Additional Topics

In addition to what will be covered in lectures and labs, there are many other NLP models you can choose from. Here are some ideas for groups looking to try something new:

- **Latent Dirichlet Allocation.** This is a generative unsupervised approach designed to group documents into topics, where topics are simply collections of similar words. In this setting, the number of topics is a hyperparameter, and the topics themselves depend on the data. This is a good method for finding themes/topics within large bodies of text. You can learn more from the original paper [here](#).
- **Language Models.** Language is often considered “self-labelled”, because even if we don’t have any other features, the sequential nature of raw text always allows us to predict the next word given the previous words in a sentence. Because of the complexity involved in predicting any possible sequence of words, these models are often high in complexity (e.g. neural nets).
- **Text Features for Non-text Prediction.** Another common approach is to use a text classification task (e.g. sentiment analysis) to derive features that can be used as input in another machine learning setting. A classic example is the use

of sentiment analysis of business news articles to predict changes in stock prices.

If you have other ideas for your analysis, feel free to discuss them with the teaching team during office hours. This list is not exhaustive, and we encourage creativity.

Part III: Presentation of Results from Part II (15 pts)

During either the last lecture (Mon, April 4th) or last lab (Wed, April 6th), groups will be required to give a brief presentation of up to 5 minutes (no longer). This presentation will cover the analysis in Part II (Part I should not be presented; include Part I results in the written report).

Following each presentation, there will be a brief 1-2 minute question period. Each group will be graded based on how well their presentation covers each of the five components described in Part II (1 pt per section, for 5 pts total).

Additionally, students will be graded on the quality of the presentation. Specifically, we will be looking for:

- Volume and pacing of speech (1 pt)
- Coherence and organization of slides (2 pts)
- Clarity of content (2 pts)
- Quality of visuals (e.g. graphs, tables) (2 pts)
- Concluding within the time limit (1 pt)
- Effectively answering questions (2 pts)

During the week of March 21, we will circulate a presentation schedule, where groups can sign up for specific time slots.